

# Virtual Memory

COMP375 Computer Architecture and  
Organization

# Cache and Android Performance

- A recent email mentioned how a programmer had to be concerned about how the cache is accessed in a multi-threaded Android program

*“Because of the cache layout, multiple threads writing to the same array risk performance degradation due to a potential high level of cache thrashing.”*

# Goals for Memory Management

- Have the programs in memory when they need to run
- Convert program addresses to physical addresses.
- Protect one program from another
- Allow programs to share data and instructions
- Support read-only or execute-only segments

# Sub Goals

- Fast address translation
- Simple for the OS to allocate and release memory
- Minimize fragmentation

This is an example of a bin packing problem and is **NP-complete** even if you have all of the memory requests at one time.

# Program to Physical

- Everybody's program is compiled and linked with the first byte of the program at address zero
- Since many programs are run simultaneously on the computer, each cannot be loaded starting at physical address zero
- The hardware converts each memory reference from a program address to a hardware physical address


# Physical Mapping Schemes

- Base Address
  - Paged Memory
  - Virtual Memory
- 
- Segmented addresses
  - Flat single address scheme

# Base Address

- Programs are located in any convenient contiguous part of memory
- When a program is to execute, the base register is loaded with the start address of the program
- The program effective address from each memory reference is added to the value in the base register to create the physical memory location

# Base Address Translation






# Limit Register

- The limit register contains the highest address of the program
- The memory system checks each memory reference to make sure it is not greater than the value in the limit register
- An addressing exception interrupt occurs if the address is too large

# Base Address with Limits




# Fragmentation

- As memory is allocated and released, the usable areas become fragmented
- **Internal fragmentation** occurs when a program is allocated to a memory area that is bigger than needed. The leftover is lost due to internal fragmentation
- **External fragmentation** occurs when there are many small unallocated areas. These areas are hard to use because they are small and divided

# Fragmented Memory

- In this memory diagram two programs use 5 MB of the 8 MB in the system
- There is 3 MB available
- A 2 MB program cannot be put in memory because the available space is fragmented



If you allocate memory in 1K blocks, what is the maximum internal fragmentation per program?

- A. 4 GB
- B. 1 KB
- C. 1023 bytes
- D. 256 bytes

# Evaluation of Base Register

- Convert program addresses to physical addresses
  - Simple and efficient
- Make it easy for the operating system to place the program in memory
  - Programs have to be located in contiguous memory. It can be a challenge to fit the programs into the available space

# Evaluation of Base Register

- Allow sharing of data and instructions
  - Each program is separate. No sharing
- Detect out of range memory accesses.
  - Simple and efficient
- Provide read-only and no-execute memory segments
  - All of memory is treated the same

# Using the Base Register Method

- Some ancient OS set the base register to certain fixed addresses to create fixed sized areas for programs. This was not very effective.
- Better OS put the programs wherever they would fit and set the base register to the beginning of the program




# Paged Memory

- RAM and programs are divided into fixed sized pages
- The page size is usually fixed for a given architecture, often between 512 -8K bytes
- The pages of a program can be put anywhere in RAM. They do not have to be contiguous.
- The page table keeps track of the physical location of pages
- The page table is indexed by the page number portion of a program address

# Pages in RAM

## RAM


### Page Table



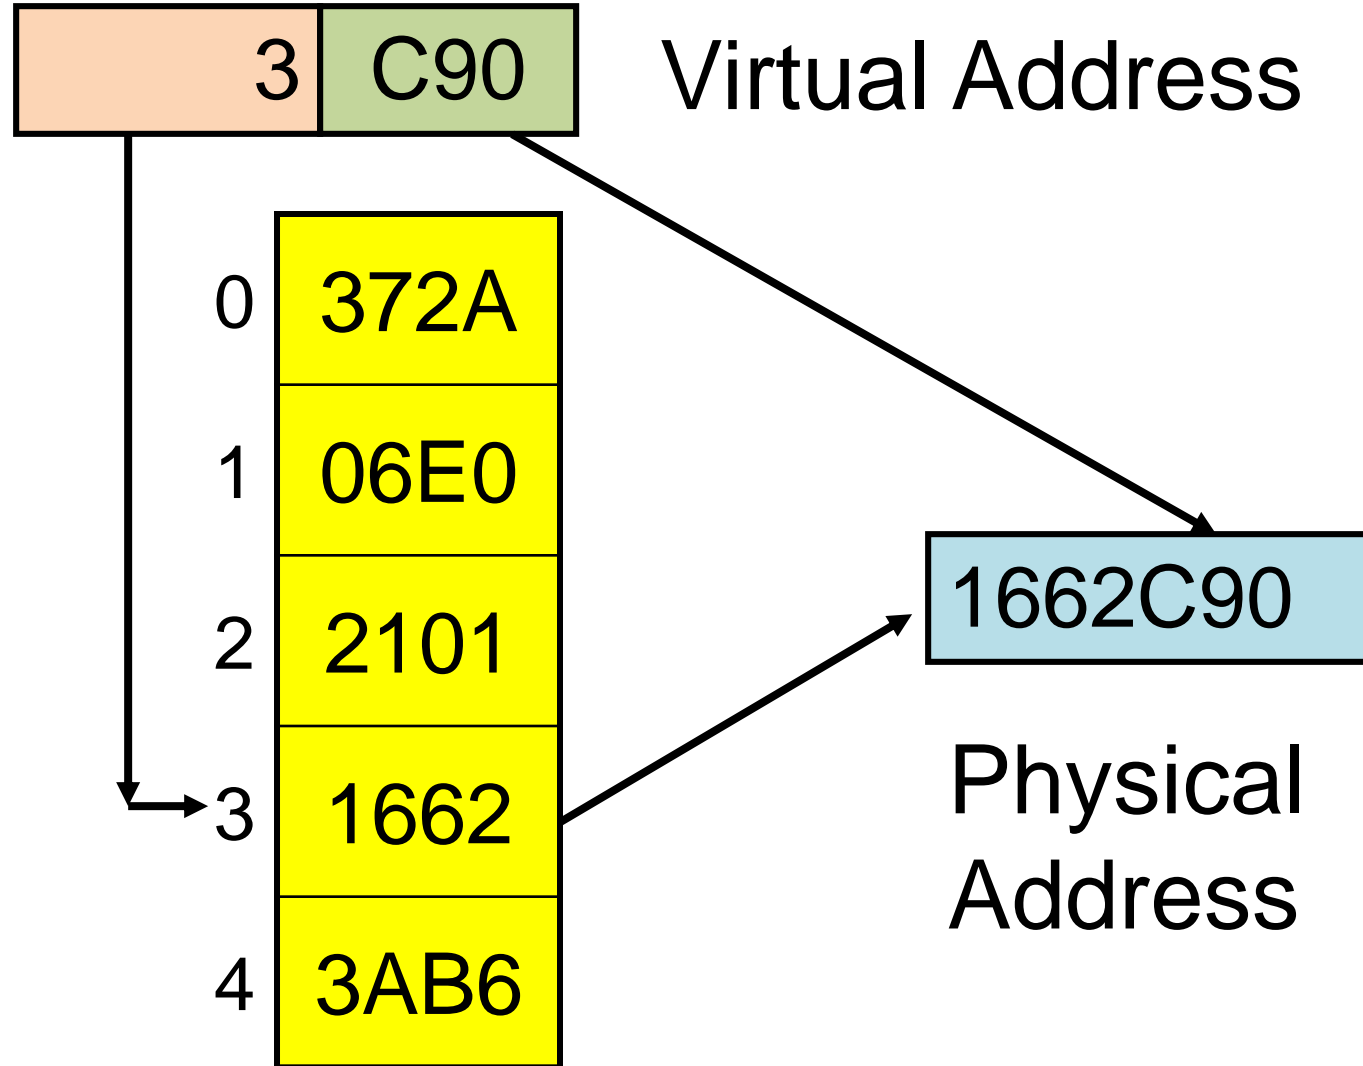
# Virtual to Physical Address

- The upper bits of a program address are used as an index into a page table
- Each page table entry contains the physical address for that page of the program
- The lower bits of the program address (which indicate which byte in the page is desired) are concatenated to the end of the physical page address

# Address Translation



## Example with 4K pages



# Dividing the Address

- The offset portion of the address is always the lower  $\log_2(\text{size of a page})$  bits of the address
- The page number portion are all of the address bits that are not part of the offset

# Fragmentation

- When using paging, programs can be loaded into RAM without worrying about fragmenting memory into small pieces. There is no **external fragmentation**.
- If the program size is not an even multiple of the page size, the last page will only be partially filled. This is **internal fragmentation**. With large memories this is unimportant.
- Paging makes much better use of RAM

How many bits are used for the  
page number and offset?

Assume **1 GB** address space and  
2K byte pages.

- A. page# = 21 offset = 11
- B. page# = 22 offset = 10
- C. page# = 19 offset = 11
- D. page# = 20 offset = 12



# Big Programs

- Even when using every byte of RAM, it is not always possible to load all the programs users would like
- Frequently large parts of programs are never executed. There are many features of Microsoft Word<sup>®</sup> you have never used.
- More programs could fit in memory if only the used portions were loaded into RAM

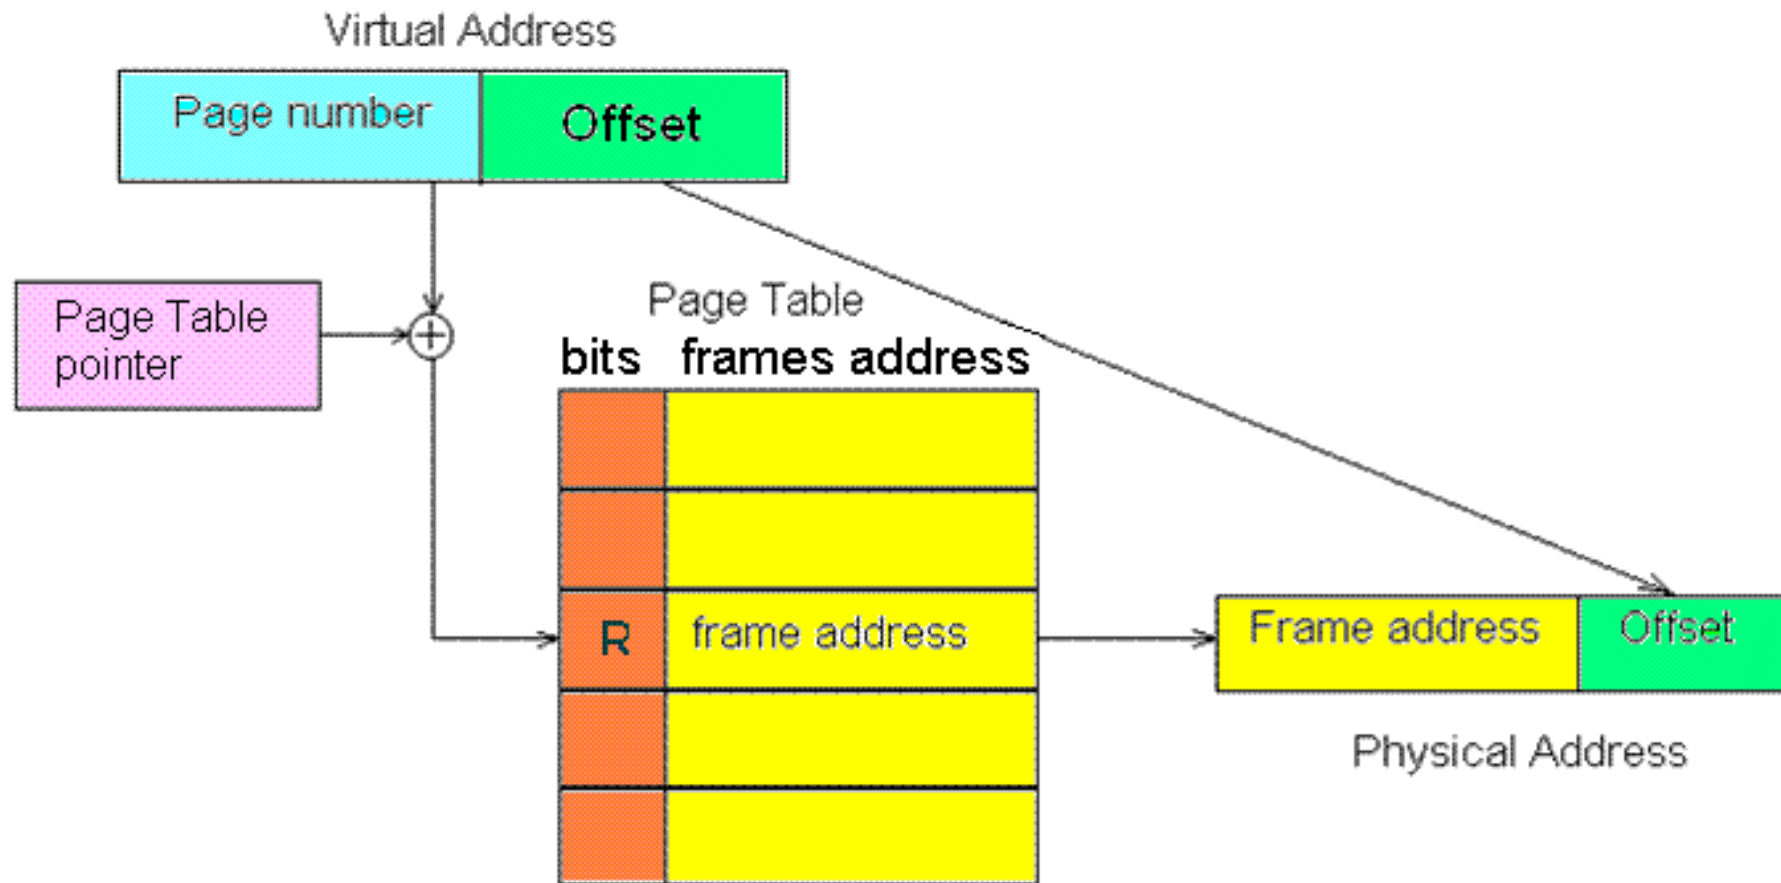
# Virtual Memory

- Virtual Memory is an extension of paging
- Only the pages that are being used are in RAM
- A copy of all pages of a program are on the page file
- If a program accesses an address in a page not in RAM, the hardware creates a page fault interrupt and the OS copies the desired page into RAM

# Virtual Memory Implementation

- Unused pages of a program do not need to be in RAM to execute the program
- A “resident” bit is added to the page table
  - Pages in RAM have the resident bit set
  - Pages not in RAM have the resident bit cleared
- All pages are stored on disk
- When a program references a page with the resident bit clear, the hardware creates a page fault interrupt

# Address Translation



# Only Necessary Pages in RAM

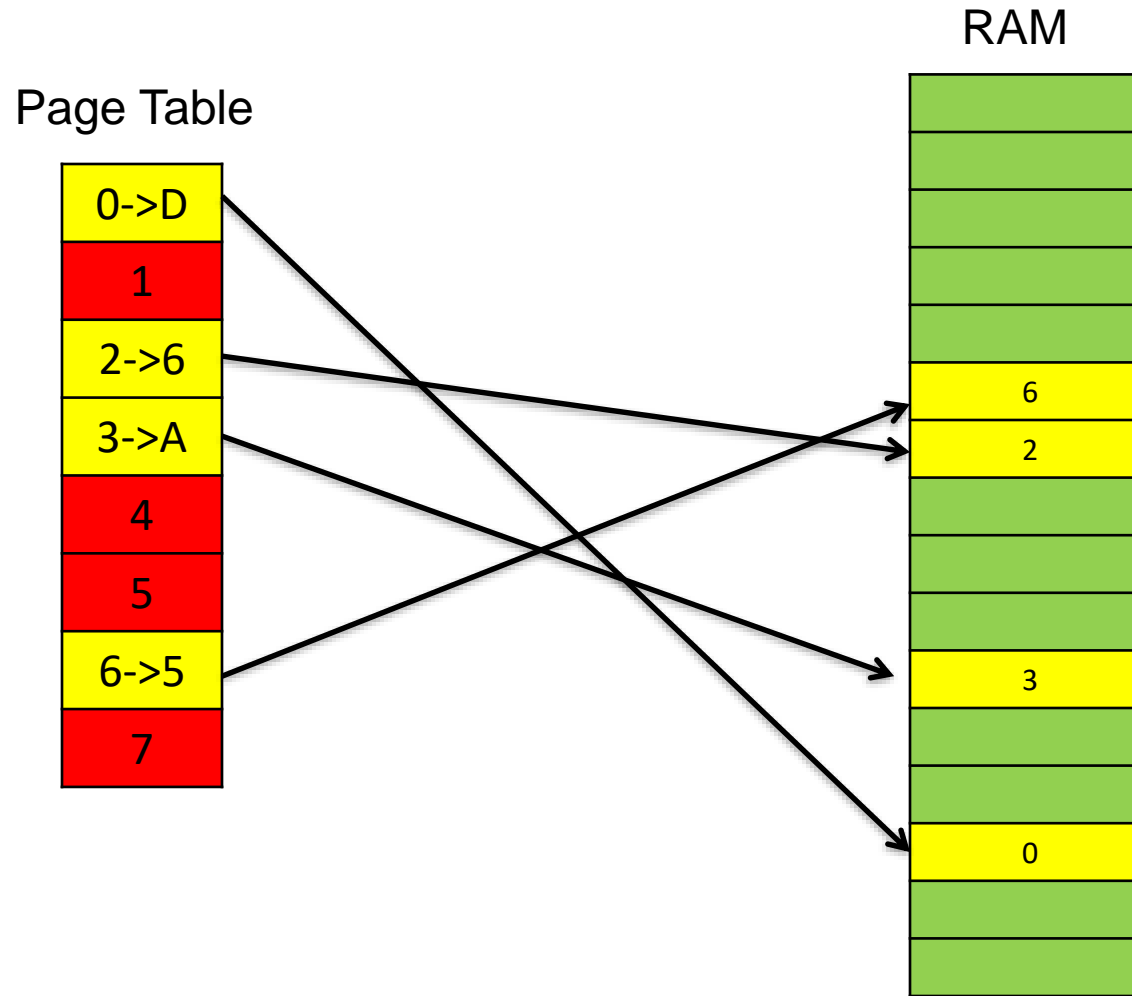
Program addresses showing pages that are referenced

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

Page Table

0->D
1
2->6
3->A
4
5
6->5
7

RAM



# Hardware and Operating System

- Virtual memory requires both hardware and operating system support
- When a page fault interrupt occurs, the OS reads the desired page from disk into an available page of RAM
- The user's page table is updated to point to the newly loaded page
- The program is placed back on the ready list to be executed

# User Programs and Virtual Memory

- Virtual memory is transparent to user programs
- Programs are unaware of page faults
- The only impact of virtual memory on user programs is performance
- If a program creates frequent page faults, the program will run very slowly

# Page Table Flags


- Each page table entry has flag bits
  - **Resident** – set if the virtual address is in RAM and clear otherwise.
  - **Used** – set when the page is referenced
  - **Dirty** – set when the page is changed
  - **No Execute** – Instructions cannot be fetched from this page
- The Used and Dirty bits are set by the hardware and cleared by the OS



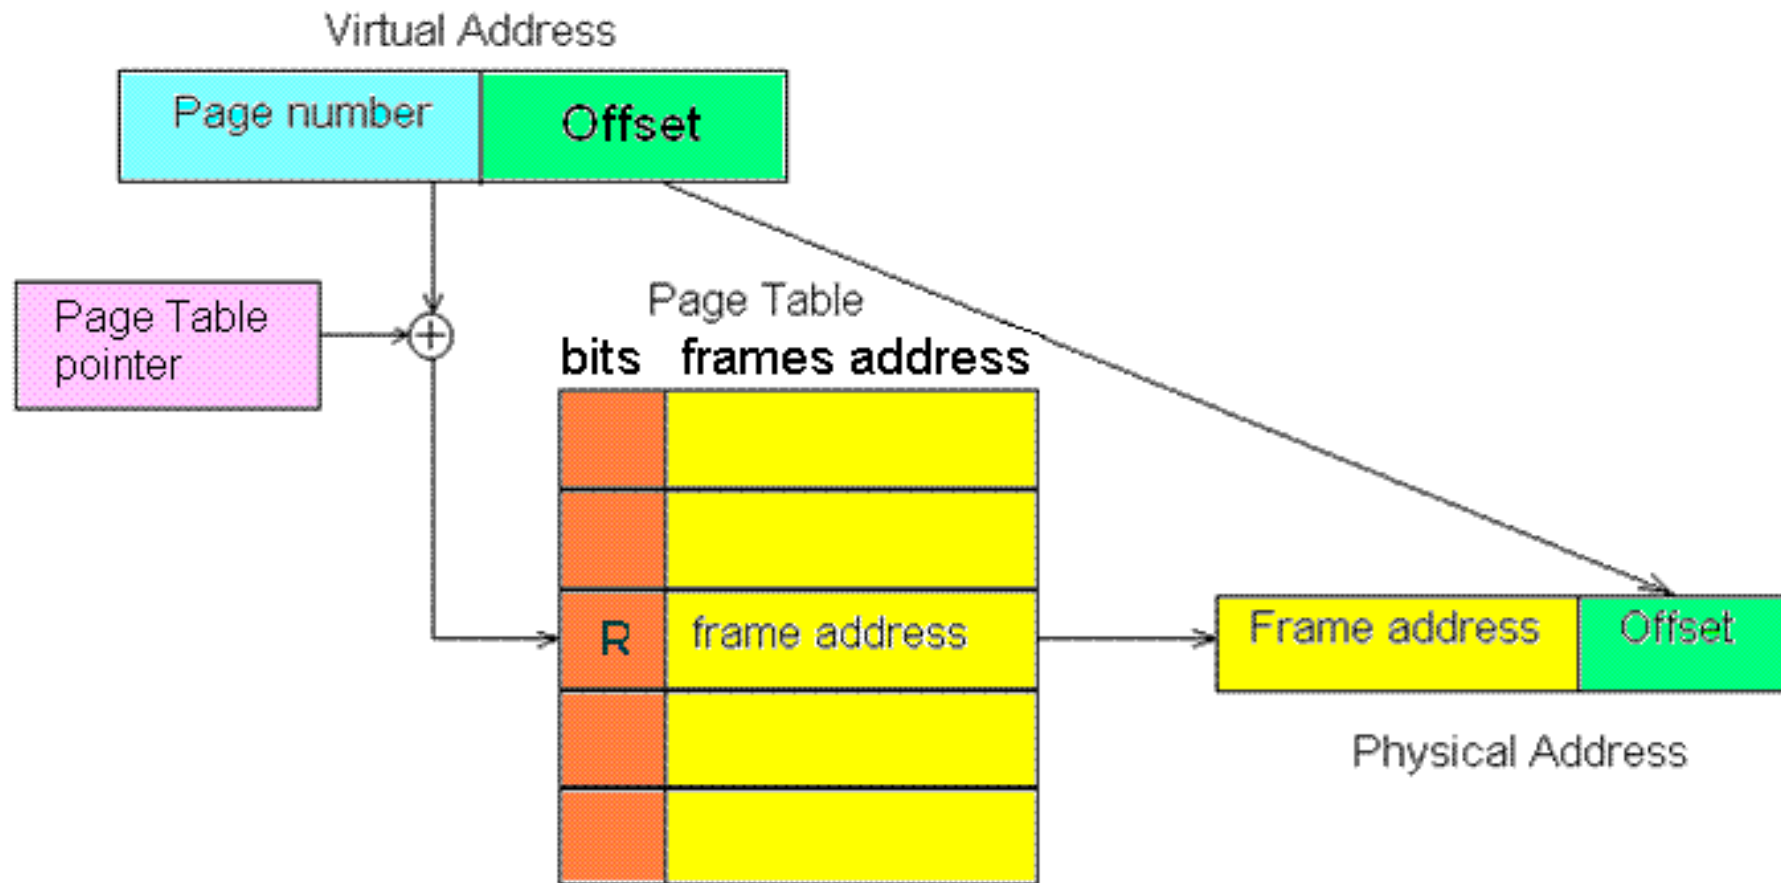
# Locality of Reference

- **Temporal locality** - a referenced location is likely to be referenced again
- **Spatial locality** - nearby locations are likely to be referenced soon

# Memory Hierarchy



# Address Translation



# Translation Lookaside Buffer

- The Translation Lookaside Buffer (TLB) is a special cache to hold recently used page table entries
- The CPU looks in the TLB for a page entry before looking in the page table
- Without the TLB each memory reference would require two memory accesses
- Most TLBs are small and use fully associative mapping

# Performance Factors

- The hardware handles the normal address translation
- When a page fault occurs, the OS determines where the page will be loaded and what page will be overwritten
- The OS must minimize the number of page faults
- The OS determines which and how many pages of a program are in RAM

# Virtual Memory Performance

- Computers can retrieve a word from RAM in about 60 ns ( $6 \times 10^{-8}$  sec)
- A good disk drive can read a block of data in about 6 ms ( $6 \times 10^{-3}$  sec)
- If the needed page is not in RAM and has to be read from the disk, it takes about 100,000 times longer to get the data.
- Page fault interrupts have to be kept to a minimum.

# A High Hit Rate Is Important

- Consider RAM in access in 60 ns and disk drive access in 6 ms
- If the hit ratio is 99.99%

$$\begin{aligned}\text{avg access} &= 0.9999 * 6 \times 10^{-8} \text{ sec} + \\ &0.0001 * 6 \times 10^{-3} \text{ sec} = 660 \text{ nsec}\end{aligned}$$

- The average memory access time is about 11 times slower

What is the average access time for  
99.9% hit rate with 9.0 ms disk  
and 9.0 ns memory?

- A. 8.1 ms
- B. 8.1 ns
- C. 9.0 ms
- D. 9.0  $\mu$ s
- E. 9.0 ns



# Program to Physical Translation Goals

- Convert program addresses to physical addresses
  - More complicated than base address. May require two memory accesses.
- Make it easy for the operating system to place the program in memory
  - Much easier than base address
- Allow sharing of data and instructions
  - Pages can be shared between programs

# Program to Physical Translation Goals

- Detect out of range memory accesses
  - If the page table index is too large, a segmentation fault interrupt will occur
- Provide read-only and no-execute memory segments
  - Different pages can be marked read-only or no-execute

# Virtual Memory Advantages

- Allows you to fit many large programs into a relatively small RAM
- Only part of a program needs to be loaded into memory
- Eliminates the need to “fit” programs into memory holes


# Virtual Memory Disadvantages

- A. Makes address translation much more complicated
- B. Can reduce performance
- C. Makes program execution time less predictable
- D. All of the above
- E. None of the above

# Segmented Memory


- Memory can be separated into segments based on the program
  - instruction segments for each function
  - data segments for global data and heap
  - stack segments
- Segmentation provides greater function isolation and makes linking easier
- Segments can have virtual memory pages

# Program Memory Organization

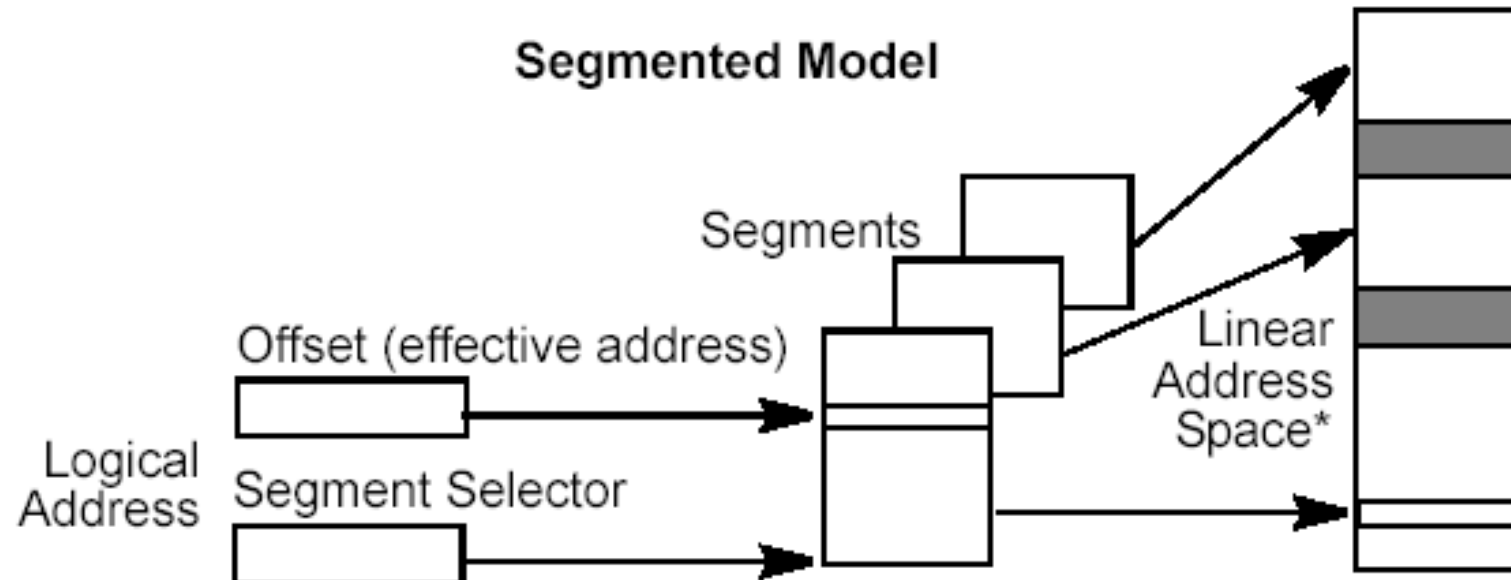


# Memory Models

## Flat Model



## Segmented Model



# Intel Segment Addresses

